


J (Theta 1)



$$h_{\theta}(x) = \theta_1 x$$

Repeat this step until achieve convergence


{  
theta sub j = theta sub j - alpha  
partial derivative J(theta's)/j

(for j = 0 & j = 1)

}

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

J (Theta 1)



$$h_{\theta}(x) = \theta_1 x$$

Repeat this step until achieve convergence


{  
theta sub j = theta sub j - alpha  
partial derivative J(theta's)/j


(for j = 0 & j = 1)

}

Gradient Descent algorithm


$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$






## Gradient Descent algorithm

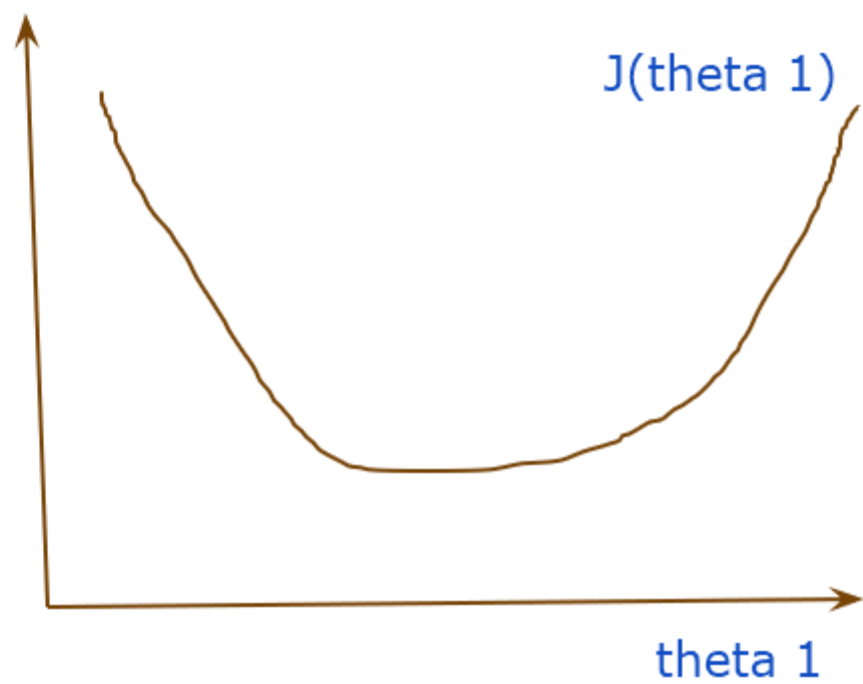
Gradient Descent algorithm



Gradient Descent algorithm




## Gradient Descent algorithm







## Gradient Descent algorithm



## Gradient Descent algorithm




## Gradient Descent algorithm



$$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$$


## Gradient Descent algorithm



$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$


$\theta_1 := \theta_1 - \alpha(\text{negative slope})$

## Gradient Descent algorithm



$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$   
 $\theta_1 := \theta_1 - \alpha(\text{negative slope})$

## Gradient Descent algorithm




$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$

$\theta_1 := \theta_1 - \alpha(\text{negative slope})$

gradient: measures how much the output of a function changes if we change the input with small values

## Gradient Descent algorithm




$$\theta_1 := \theta_1 - \alpha (\text{partial derivative } J(\theta_1) / \theta_1)$$

$$\theta_1 := \theta_1 - \alpha (\text{negative slope})$$

gradient: measures how much the output of a function changes if we change the input with small values

## Gradient Descent algorithm




$$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$$

$$\theta_1 := \theta_1 - \alpha (\text{negative slope})$$

gradient: measures how much the output of a function changes if we change the input with small values



## Gradient Descent algorithm




$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$

$\theta_1 := \theta_1 - \alpha(\text{negative slope})$

gradient: measures how much the output of a function changes if we change the input with small values

## Gradient Descent algorithm




$$\theta_1 := \theta_1 - \alpha (\text{partial derivative } J(\theta_1)/\theta_1)$$

$$\theta_1 := \theta_1 - \alpha (\text{negative slope})$$

gradient: measures how much the output of a function changes if we change the input with small values

## Gradient Descent algorithm




$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$

$\theta_1 := \theta_1 - \alpha$  (negative slope)

gradient: measures how much the output of a function changes if we change the input with small values

## Gradient Descent algorithm




$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$

$\theta_1 := \theta_1 - \alpha$  (negative slope)

gradient: measures how much the output of a function changes if we change the input with small values

## Gradient Descent algorithm




$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$

$\theta_1 := \theta_1 - \alpha$  (negative slope)

gradient: measures how much the output of a function changes if we change the input with small values

## Gradient Descent algorithm




$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$

$\theta_1 := \theta_1 - \alpha$  (negative slope)

gradient: measures how much the output of a function changes if we change the input with small values

## Gradient Descent algorithm




$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$

$\theta_1 := \theta_1 - \alpha$  (negative slope)

gradient: measures how much the output of a function changes if we change the input with small values

## Gradient Descent algorithm




$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$

$\theta_1 := \theta_1 - \alpha$  (negative slope)

gradient: measures how much the output of a function changes if we change the input with small values



## Gradient Descent algorithm




$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$

$\theta_1 := \theta_1 - \alpha$  (negative slope)

gradient: measures how much the output of a function changes if we change the input with small values

## Gradient Descent algorithm




$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$

$\theta_1 := \theta_1 - \alpha$  (negative slope)

gradient: measures how much the output of a function changes if we change the input with small values

## Gradient Descent algorithm




$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$

$\theta_1 := \theta_1 - \alpha$  (negative slope)

gradient: measures how much the output of a function changes if we change the input with small values

## Gradient Descent algorithm




$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$

$\theta_1 := \theta_1 - \alpha$  (negative slope)

gradient: measures how much the output of a function changes if we change the input with small values

## Gradient Descent algorithm



$\theta_1 := \theta_1 - \alpha(\text{partial derivative } J(\theta_1)/\theta_1)$

$\theta_1 := \theta_1 - \alpha$  (negative slope)

gradient: measures how much the output of a function changes if we change the input with small values

# Gradient Descent Algorithm

Have some function  $J(\theta_0, \theta_1)$

Want  $\min_{\theta_0, \theta_1} J(\theta_0, \theta_1)$  This is true for all  
For  $i = 0, 1, 2, 3, \dots, n$   $\theta_i$

## Steps:

- Start with some  $\theta_0, \theta_1$
- Keep changing  $\theta_0, \theta_1$  to reduce  $J(\theta_0, \theta_1)$   
until we hopefully end up at a minimum

# Gradient descent algorithm

repeat until convergence {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (\text{for } j = 0 \text{ and } j = 1)$$

}

  
**Learning Rate**

---

Correct Update:

$$\text{temp0} := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

$$\text{temp1} := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_0 := \text{temp0}$$

$$\theta_1 := \text{temp1}$$

update  
Simultaneously:  $\theta_0$  and  $\theta_1$

Gradient descent can converge to a local minimum, even with fixed ( $\alpha$ ) learning rate.

$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$$

As we approach a local minimum, gradient descent will automatically take smaller steps. So, no need to decrease  $\alpha$  over time.



"Batch" gradient Descent algorithm:  
"Stochastic" gradient descent algorithm

"Batch" gradient Descent algorithm:  
"Stochastic" gradient descent algorithm

Each step of gradient descent uses all the training example: batch gradient descent algorithm

Each step of gradient decent uses single training example

Mini-batch: each step of gradient descent uses subset of training example

"Batch" gradient Descent algorithm:  
"Stochastic" gradient descent algorithm

Each step of gradient descent uses all the training example: batch gradient descent algorithm

Each step of gradient decent uses single training example

Mini-batch: each step of gradient descent uses subset of training example

## **“Batch” Gradient Descent**

“Batch”: Each step of gradient descent uses all the training examples.

## **“Stochastic” Gradient Descent**

“Stochastic”: Each step of gradient descent use only single training example.

## **“Mini-batch” Gradient Descent**

It lies in between of these two extremes, and can use a mini-batch (small portion) of training data examples in each step

## Batch Gradient Descent Algorithm

## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \sum_{i=1}^m [h(x)_i - (y)_i]^2$$

⌋

## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \sum_{i=1}^m [h(x)^i - (y)^i]^2$$

$\theta_1$ :  $\theta_1 - \alpha$  (partial derivative of cost function with respect to the parameters)

## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \text{summation } (i = 1 \text{ to } m) [h(x)_i - (y)_i]^2$$

$\theta_1$ :  $\theta_1 - (\alpha) (\text{partial derivative of cost function with respect to the parameters})$



## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \text{summation } (i = 1 \text{ to } m) [h(x)_i - (y)_i]^2$$

$\theta_1$ :  $\theta_1 - (\alpha) (\text{partial derivative of cost function with respect to the parameters})$



## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \text{summation } (i = 1 \text{ to } m) [h(x)_i - (y)_i]^2$$

$\theta_1$ :  $\theta_1 - (\alpha) (\text{partial derivative of cost function with respect to the parameters})$

Training data points

## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \text{summation } (i = 1 \text{ to } m) [h(x)_i - (y)_i]^2$$

$\theta_1$ :  $\theta_1 - (\alpha) (\text{partial derivative of cost function with respect to the parameters})$

Training data points

Vanilla gradient descent algorithm/gradient descent algorithm

## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \text{summation } (i = 1 \text{ to } m) [h(x)_i - (y)_i]^2$$

$\theta_1$ :  $\theta_1 - (\alpha) (\text{partial derivative of cost function with respect to the parameters})$

Training data points

Vanilla gradient descent algorithm/gradient descent algorithm

Stochastic Gradient descent algorithm

## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \text{summation } (i = 1 \text{ to } m) [h(x)_i - (y)_i]^2$$

$\theta_1$ :  $\theta_1 - (\alpha) (\text{partial derivative of cost function with respect to the parameters})$

Training data points

Vanilla gradient descent algorithm/gradient descent algorithm

Stochastic Gradient descent algorithm

update the parameters for each training example one by one

## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \text{summation } (i = 1 \text{ to } m) [h(x)_i - (y)_i]^2$$

$\theta_1$ :  $\theta_1 - (\alpha) (\text{partial derivative of cost function with respect to the parameters})$



Training data points

Vanilla gradient descent algorithm/gradient descent algorithm

Stochastic Gradient descent algorithm

update the parameters for each training example one by one

Mini-batch gradient descent algorithm

## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \text{summation } (i = 1 \text{ to } m) [h(x)_i - (y)_i]^2$$

$\theta_1$ :  $\theta_1 - (\alpha) (\text{partial derivative of cost function with respect to the parameters})$



Training data points

Vanilla gradient descent algorithm/gradient descent algorithm

Stochastic Gradient descent algorithm

update the parameters for each training example one by one

Mini-batch gradient descent algorithm : splits the training dataset into small batches & perform an update

## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \text{summation } (i = 1 \text{ to } m) [h(x)_i - (y)_i]^2$$

$\theta_1$ :  $\theta_1 - (\alpha) (\text{partial derivative of cost function with respect to the parameters})$

Training data points

Vanilla gradient descent algorithm/gradient descent algorithm

Stochastic Gradient descent algorithm

update the parameters for each training example one by one

Mini-batch gradient descent algorithm : splits the training dataset into small batches & perform an update




## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \text{summation } (i = 1 \text{ to } m) [h(x)_i - (y)_i]^2$$

$\theta_1$ :  $\theta_1 - (\alpha) (\text{partial derivative of cost function with respect to the parameters})$

↓  
Training data points



Vanilla gradient descent algorithm/gradient descent algorithm

Stochastic Gradient descent algorithm

update the parameters for each training example one by one


Mini-batch gradient descent algorithm : splits the training dataset into small batches & perform an update

## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \text{summation } (i = 1 \text{ to } m) [h(x)_i - (y)_i]^2$$

$\theta_1$ :  $\theta_1 - (\alpha) (\text{partial derivative of cost function with respect to the parameters})$

↓  
Training data points



Vanilla gradient descent algorithm/gradient descent algorithm

Stochastic Gradient descent algorithm

update the parameters for each training example one by one


Mini-batch gradient descent algorithm : splits the training dataset into small batches & perform an update

## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \text{summation } (i = 1 \text{ to } m) [h(x)_i - (y)_i]^2$$

$\theta_1$ :  $\theta_1 - (\alpha) (\text{partial derivative of cost function with respect to the parameters})$

Training data points



Vanilla gradient descent algorithm/gradient descent algorithm

Stochastic Gradient descent algorithm

update the parameters for each training example one by one


Mini-batch gradient descent algorithm : splits the training dataset into small batches & perform an update

## Batch Gradient Descent Algorithm

$$J(\theta_1) := \frac{1}{2m} \text{summation } (i = 1 \text{ to } m) [h(x)_i - (y)_i]^2$$

$\theta_1$ :  $\theta_1 - (\alpha) (\text{partial derivative of cost function with respect to the parameters})$

↓  
Training data points



Vanilla gradient descent algorithm/gradient descent algorithm

Stochastic Gradient descent algorithm

update the parameters for each training example one by one

Mini-batch gradient descent algorithm : splits the training dataset into small batches & perform an update